

Design Notes Assignment 3

Majid Ghaderi

Disclaimer

These notes are based on my own implementation. I do not claim that my implementation is the simplest or the best. Feel free to use or disregard any of these suggestions as you wish.

General Tips

- Start with basic UDP socket IO. Implement a simple client/server program to send and receive Segments using UDP.
- Start testing with small files that have just one segment (*e.g.*, file size is 1 byte).
- Add a lot of debug messages to your code. You can remove/disable them later when not needed. Do not reinvent the wheel here. Look at the Java debugging facilities specifically the Logger class.

Program Structure

The high-level structure of my `send()` method is as follows:

-
- 1: complete TCP handshake
 - 2: start ACK receiving thread
 - 3: **while** not end of file **do**
 - 4: create a segment `seg` with next sequence number
 - 5: **while** (`txQueue.isFull()`) wait
 - 6: `processSend(seg)`
 - 7: **end while**
 - 8: **while** (`not txQueue.isEmpty()`) wait
 - 9: send end of transmission to server over TCP
 - 10: cancel timer, ACK receiving thread, clean up
-

Rather than simply *waiting* for the txQueue to become empty or not full, which wastes CPU time, a better strategy is to yield the execution so that other threads can run by calling Thread.yield().

I use a separate thread to receive ACKs and a timer task to schedule time-outs. I have also defined three helper methods to handle sending segments, receiving ACKs and retransmissions due to time-out:

```
public synchronized void processSend(Segment seg) {
    // send seg to the UDP socket
    // add seg to the transmission queue txQueue
    // if txQueue.size() == 1, start the timer
}

public synchronized void processACK(Segment ack) {
    // if ACK not in the current window, do nothing
    // otherwise:
    // cancel the timer
    // while txQueue.element().getSeqNum() < ack.getSeqNum()
    //     txQueue.remove()
    // if not txQueue.isEmpty(), start the timer
}

public synchronized void processTimeout() {
    // get the list of all pending segments by calling txQueue.toArray()
    // go through the list and send all segments to the UDP socket
    // if not txQueue.isEmpty(), start the timer
}
```

The methods processACK() and processTimeout() are called by the ACK receiving thread and the timer task, respectively, as explained below. As you can see, these three methods are synchronized to avoid running them simultaneously, which could results in unintended concurrency problems as they all modify the content of txQueue.

Setting up a Timer

Class Timer can be used to schedule a recurring timer. To use the Timer class, you need to define a timer task class as well. A timer task is similar to a Thread class with a run() method. The only difference is that it extends class TimerTask. Here is an example of a timer task class:

```
class TimeoutHandler extends TimerTask {  
  
    // define constructor  
  
    // the run method  
    public void run() {  
        // call processTimeout() in the main class  
    }  
}
```

The following piece of code demonstrates how to start a timer that goes off in 1000 milli-seconds,

```
Timer timer = new Timer(true);  
timer.schedule(new TimeoutHandler(), 1000);
```

A timer can be cancelled by calling its `cancel()` method.

Receiving ACKs

Since arrival of ACK segments happens in parallel with sending segments, a separate thread can be started to receive ACKs from the same UDP socket that is used for sending segments.

```
public class ReceiverThread extends Thread {  
  
    // define constructor  
  
    // the run method  
    public void run() {  
        // while not terminated:  
        // 1. receive a DatagramPacket pkt from UDP socket  
        // 2. call processAck(new Segment(pkt)) in the parent process  
    }  
}
```